Floating Point

COE 301 / ICS 233

Computer Organization

Dr. Muhamed Mudawar

College of Computer Sciences and Engineering King Fahd University of Petroleum and Minerals

Presentation Outline

Floating-Point Numbers

The IEEE 754 Floating-Point Standard

- Floating-Point Comparison, Addition and Subtraction
- Floating-Point Multiplication
- MIPS Floating-Point Instructions and Examples

The World is Not Just Integers

- Programming languages support numbers with fraction
 - Called floating-point numbers
 - **Examples**:

```
3.14159265...(\pi)
```

2.71828... (*e*)

 1.0×10^{-9} (seconds in a nanosecond)

 8.64×10^{13} (nanoseconds in a day)

The last number is a large integer that cannot fit in a 32-bit register

- * We use a scientific notation to represent
 - \diamond Very small numbers (e.g. 1.0×10^{-9})
 - \diamond Very large numbers (e.g. 8.64 × 10¹³)
 - ♦ Scientific notation: ± d. fraction × 10 ± exponent

Floating-Point Numbers

Examples of floating-point numbers in base 10

Examples of floating-point numbers in base 2

$$-1.00101 \times 2^{23}$$
 , 1.101101×2^{-3} binary point

- Exponents are kept in decimal for clarity
- Floating-point numbers should be normalized
 - Exactly one non-zero digit should appear before the point
 - In a decimal number, this digit can be from 1 to 9
 - In a binary number, this digit should be 1
 - ◇ Normalized: -5.341×10³ and 1.101101×2⁻³
 - ◆ NOT Normalized: -0.05341×10⁵ and 1101.101×2⁻⁶

Floating-Point Representation

A floating-point number is represented by the triple

- Sign bit (0 is positive and 1 is negative)
 - Representation is called sign and magnitude
- Exponent field (signed value)
 - Very large numbers have large positive exponents
 - Very small close-to-zero numbers have negative exponents
 - More bits in exponent field increases range of values
- Fraction field (fraction after binary point)
 - More bits in fraction field improves the precision of FP numbers



IEEE 754 Floating-Point Standard

- Found in virtually every computer invented since 1980
 - Simplified porting of floating-point numbers
 - Unified the development of floating-point algorithms
 - Increased the accuracy of floating-point numbers
- Single Precision Floating Point Numbers (32 bits)
 - ♦ 1-bit sign + 8-bit exponent + 23-bit fraction

S	Exponent ⁸	Fraction ²³
---	-----------------------	------------------------

- Double Precision Floating Point Numbers (64 bits)
 - ♦ 1-bit sign + 11-bit exponent + 52-bit fraction

S	Exponent ¹¹	Fraction ⁵²								
(continued)										

Normalized Floating Point Numbers

For a normalized floating point number (S, E, F)

$$F = f_1 f_2 f_3 f_4 \dots$$

- Significand is equal to $(1.F)_2 = (1.f_1f_2f_3f_4...)_2$
 - ♦ IEEE 754 assumes hidden 1. (not stored) for normalized numbers
 - ♦ Significand is 1 bit longer than fraction
- Value of a Normalized Floating Point Number:

$$\pm (\mathbf{1}.F)_{2} \times 2^{\text{exponent_value}}$$

$$\pm (\mathbf{1}.f_{1}f_{2}f_{3}f_{4}...)_{2} \times 2^{\text{exponent_value}}$$

$$\pm (\mathbf{1} + f_{1} \times 2^{-1} + f_{2} \times 2^{-2} + f_{3} \times 2^{-3} + f_{4} \times 2^{-4}...)_{2} \times 2^{\text{exponent_value}}$$

$$S = 0$$
 is positive, $S = 1$ is negative

Biased Exponent Representation

- * How to represent a signed exponent? Choices are ...
 - ♦ Sign + magnitude representation for the exponent
 - ♦ Two's complement representation
 - Biased representation
- ❖ IEEE 754 uses biased representation for the exponent
 - \diamond Exponent Value = E Bias (Bias is a constant)
- The exponent field is 8 bits for single precision
 - \diamond E can be in the range 0 to 255
 - ightharpoonup E = 0 and E = 255 are reserved for special use (discussed later)
 - ightharpoonup E = 1 to 254 are used for normalized floating point numbers
 - ♦ Bias = 127 (half of 254)
 - \diamond Exponent value = E 127 Range: -126 to +127

Biased Exponent – Cont'd

- For double precision, the exponent field is 11 bits
 - \diamond *E* can be in the range 0 to 2047
 - \Rightarrow E = 0 and E = 2047 are reserved for special use
 - ightharpoonup E = 1 to 2046 are used for normalized floating point numbers
 - ♦ Bias = 1023 (half of 2046)
 - ♦ Exponent value = E 1023 Range: -1022 to +1023
- Value of a Normalized Floating Point Number is

$$\pm (1.F)_{2} \times 2^{(E-Bias)}$$

$$\pm (1.f_{1}f_{2}f_{3}f_{4}...)_{2} \times 2^{(E-Bias)}$$

$$\pm (1 + f_{1} \times 2^{-1} + f_{2} \times 2^{-2} + f_{3} \times 2^{-3} + f_{4} \times 2^{-4} ...)_{2} \times 2^{(E-Bias)}$$

$$S = 0$$
 is positive, $S = 1$ is negative

Examples of Single Precision Float

What is the decimal value of this Single Precision float?

Solution:

- ♦ Sign = 1 is negative
- \Rightarrow E = (01111100)₂ = 124, E bias = 124 127 = -3
- \Rightarrow Significand = $(1.0100 \dots 0)_2 = 1 + 2^{-2} = 1.25$ (1. is implicit)
- ♦ Value in decimal = $-1.25 \times 2^{-3} = -0.15625$
- What is the decimal value of?
- Solution: implicit
 - \diamond Value in decimal = +(1.01001100 ... 0)₂ × 2¹³⁰⁻¹²⁷ = (1.01001100 ... 0)₂ × 2³ = (1010.01100 ... 0)₂ = 10.375

Examples of Double Precision Float

* What is the decimal value of this Double Precision float?

0	1	- 0	0	0	0	0	0	0	1	0	1	0	0	1	0	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

Solution:

- \diamond Value of exponent = $(10000000101)_2$ Bias = 1029 1023 = 6
- ightharpoonup Value of double = $(1.00101010 \dots 0)_2 \times 2^6$ (1. is implicit) = $(1001010.10 \dots 0)_2 = 74.5$
- What is the decimal value of?

• Do it yourself! (answer should be $-1.5 \times 2^{-7} = -0.01171875$)

Decimal to Binary Floating-Point

- Convert –0.8125 to single and double-precision floating-point
- Solution:
 - Fraction bits can be obtained using multiplication by 2

■
$$0.8125 \times 2$$
 = 1.625
■ 0.625×2 = 1.25 $0.8125 = (0.1101)_2 = \frac{1}{2} + \frac{1}{4} + \frac{1}{16} = 13/16$
■ 0.5×2 = 1.0

- Stop when fractional part is 0, or after computing all required fraction bits
- ♦ Fraction = $(0.1101)_2$ = $(1.101)_2 \times 2^{\frac{1}{2}}$ (Normalized)
- \Rightarrow Exponent = (-1) + Bias = 126 (single precision) and 1022 (double)

1	0	1	1	1	1	1	1	0	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	Single Precision
1	0	1	1	1	1	1	1	1	1	1	0	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	Double
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	Precision

Largest Normalized Float

- What is the Largest normalized float?
- Solution for Single Precision:

- \triangle E bias = 254 127 = +127 (largest exponent for SP)
- \Rightarrow Significand = $(1.111 ... 1)_2$ = 1.99999988 = almost 2
- ♦ Value in decimal $\approx 2 \times 2^{+127} \approx 2^{+128} \approx 3.4028 \dots \times 10^{+38}$
- Solution for Double Precision:

- ♦ Value in decimal $\approx 2 \times 2^{+1023} \approx 2^{+1024} \approx 1.79769 \dots \times 10^{+308}$
- Overflow: exponent is too large to fit in the exponent field

Smallest Normalized Float

- What is the smallest (in absolute value) normalized float?
- Solution for Single Precision:

 - \diamond Exponent bias = 1 127 = -126 (smallest exponent for SP)
 - \Rightarrow Significand = $(1.000 ... 0)_2 = 1$
 - \diamond Value in decimal = 1 × 2⁻¹²⁶ = 1.17549 ... × 10⁻³⁸
- Solution for Double Precision:
 - 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 0
 - \diamond Value in decimal = 1 × 2⁻¹⁰²² = 2.22507 ... × 10⁻³⁰⁸
- Underflow: exponent is too small to fit in exponent field

Zero, Infinity, and NaN

Zero

- \Rightarrow Exponent field E = 0 and fraction F = 0
- ♦ +0 and -0 are both possible according to sign bit S

Infinity

- ightharpoonup Infinity is a special value represented with maximum E and F = 0
 - For single precision with 8-bit exponent: maximum E = 255
 - For double precision with 11-bit exponent: maximum E = 2047
- Infinity can result from overflow or division by zero
- \diamond + ∞ and - ∞ are both possible according to sign bit S

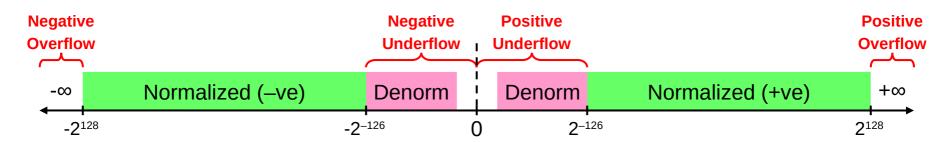
* NaN (Not a Number)

- \diamond NaN is a special value represented with maximum E and $F \neq 0$
- \diamond 0 / 0 \clubsuit NaN, 0 × ∞ \clubsuit NaN, sqrt(-1) \clubsuit NaN
- Operation on a NaN is typically a NaN: Op(X, NaN) & NaN

Denormalized Numbers

- LEEE standard uses denormalized numbers to ...
 - ♦ Fill the gap between 0 and the smallest normalized float
 - Provide gradual underflow to zero
- Denormalized: exponent field E is 0 and fraction $F \neq 0$
 - ♦ The Implicit 1. before the fraction now becomes 0. (denormalized)
- \diamond Value of denormalized number (S, 0, F)

Single precision: $\pm (0.F)_2 \times 2^{-126}$ Double precision: $\pm (0.F)_2 \times 2^{-1022}$



Summary of IEEE 754 Encoding

Single-Precision	Exponent = 8	Fraction = 23	Value				
Normalized Number	1 to 254	Anything	$\pm (1.F)_2 \times 2^{E-127}$				
Denormalized Number	0	nonzero	$\pm (0.F)_2 \times 2^{-126}$				
Zero	0	0	± 0				
Infinity	255	0	± ∞				
NaN	255	nonzero	NaN				

Double-Precision	Exponent = 11	Fraction = 52	Value			
Normalized Number	1 to 2046	Anything	$\pm (1.F)_2 \times 2^{E-1023}$			
Denormalized Number	0	nonzero	$\pm (0.F)_2 \times 2^{-1022}$			
Zero	0	0	± 0			
Infinity	2047	0	± ∞			
NaN	2047	nonzero	NaN			



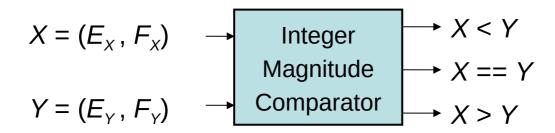
- Floating-Point Numbers
- The IEEE 754 Floating-Point Standard
- * Floating-Point Comparison, Addition and Subtraction
- Floating-Point Multiplication
- MIPS Floating-Point Instructions and Examples

Floating-Point Comparison

- IEEE 754 floating point numbers are ordered (except NaN)
 - ♦ Because the exponent uses a biased representation ...
 - Exponent value and its binary representation have same ordering
 - Placing exponent before the fraction field orders the magnitude
 - Larger exponent ⇒ larger magnitude
 - For equal exponents, Larger fraction ⇒ larger magnitude

■
$$0 < (0.F)_2 \times 2^{Emin} < (1.F)_2 \times 2^{E-Bias} < \infty$$
 $(E_{min} = 1 - Bias)$

- Sign bit provides a quick test for signed <</p>
- Integer comparator can compare the magnitudes



Floating Point Addition

Consider Adding Single-Precision Floats:

- $1.1000000000000110000101_2 \times 2^2$
- Cannot add significands ... Why?
 - Because exponents are not equal
- * How to make exponents equal?
 - Shift the significand of the lesser exponent right
 - \diamond Difference between the two exponents = 4 2 = 2
 - So, shift right second number by 2 bits and increment exponent

$$1.1000000000000110000101_2 \times 2^2$$

 $= 0.0110000000000001100001 01, \times 2^{4}$

Floating-Point Addition — cont'd

Now, ADD the Significands:

```
\times 2^4
```

$$+ 1.1000000000000110000101 \times 2^{2}$$

- + 0.01100000000000001100001 01 × 24 (shift right)
 - **10.0100010000000001100011 01** \times 2⁴ (result)
- * Addition produces a carry bit, result is NOT normalized
- Normalize Result (shift right and increment exponent):

```
10.0100010000000001100011 01 \times 24
```

 $= 1.00100010000000000110001 101 \times 2^{5}$ (normalized)

Rounding

- Single-precision requires only 23 fraction bits
- However, Normalized result can contain additional bits

1.0010001000000000110001 |
$$(1)(01) \times 2^5$$

Round Bit: R = 1 $\xrightarrow{1}$ Sticky Bit: S = 1

- Two extra bits are used for rounding
 - Round bit: appears just after the normalized result
 - Sticky bit: appears after the round bit (OR of all additional bits)
- Since **RS = 11**, increment fraction to **round to nearest**
 - $1.0010001000000000110001 \times 2^{5}$

+1

Floating-Point Subtraction

- Addition is used when operands have the same sign
- Addition becomes a subtraction when sign bits are different
- Consider adding floating-point numbers with different signs:

```
+ 1.00000000101100010001101 \times 2<sup>-6</sup>
```

- $-1.00000000000000010011010 \times 2^{-1}$
- + 0.00001000000001011000100 01101×2^{-1} (shift right 5 bits)
- -1.00000000000000010011010 $\times 2^{-1}$
- $0.00001000000001011000100 01101 \times 2^{-1}$
- 1 0.11111111111111101100110 \times 2⁻¹ (2's complement)
- 1 1.00001000000001000101010 01101 \times 2⁻¹ (Negative result)
- 0.111101111111110111010101 10011 × 2⁻¹ (Sign Magnitude)
- 2's complement of result is required if result is negative

Floating-Point Subtraction — cont'd

- + 1.00000000101100010001101 \times 2⁻⁶
- $-1.00000000000000010011010 \times 2^{-1}$
- 0.111101111111110111010101 10011 × 2⁻¹ (Sign Magnitude)
- * Result should be normalized (unless it is equal to zero)
 - For subtraction, we can have leading zeros. To normalize, count the number of leading zeros, then shift result left and decrement the exponent accordingly.

Guard bit

- 0.111101111111101110010101 1 0011 × 2⁻¹
- 1.11101111111110111010101011 0011 × 2⁻² (Normalized)
- Guard bit: guards against loss of a fraction bit
 - Needed for subtraction only, when result has a leading zero and should be normalized.

Floating-Point Subtraction — cont'd

Next, the normalized result should be rounded

Guard bit

- 0.11110111111110111001010101 1 0 011 × 2⁻¹
- 1.1110111111111110111010101011()([0])011 × 2⁻² (Normalized) Round bit: R=0 -- -- Sticky bit: S = 1
- ❖ Since R = 0, it is more accurate to truncate the result even though S = 1. We simply discard the extra bits.
 - 1.11101111111101110101011 0 011 \times 2⁻² (Normalized)
 - 1.11101111111101110101011 \times 2⁻² (Rounded to nearest)
- IEEE 754 Representation of Result

Rounding to Nearest Even

- Normalized result has the form: **1.** $f_1 f_2 \dots f_l R S$
 - \diamond The round bit **R** appears immediately after the last fraction bit f_i
 - ♦ The sticky bit S is the OR of all remaining additional bits
- Round to Nearest Even: default rounding mode
- * Four cases for RS:
 - \diamond RS = 00 \clubsuit Result is Exact, no need for rounding
 - ♦ RS = 01 ♣ Truncate result by discarding RS
 - RS = 11 A Increment result: ADD 1 to last fraction bit
 - \diamond RS = 10 \clubsuit Tie Case (either truncate or increment result)
 - Check Last fraction bit f_1 (f_{23} for single-precision or f_{52} for double)
 - If f_i is 0 then truncate result to keep fraction even
 - If f, is 1 then increment result to make fraction even

Additional Rounding Modes

- LEEE 754 standard includes other rounding modes:
- 1. Round to Nearest Even: described in previous slide
- Round toward +Infinity: result is rounded up
 Increment result if sign is positive and R or S = 1
- 3. Round toward -Infinity: result is rounded down Increment result if sign is negative and R or S = 1
- 4. Round toward 0: always truncate result
- * Rounding or Incrementing result might generate a carry
 - This occurs only when all fraction bits are 1
 - ♦ Re-Normalize after Rounding step is required only in this case

Example on Rounding

- Round following result using IEEE 754 rounding modes:
- Round to Nearest Even: Round Bit Sticky Bit
 - \diamond Increment result since RS = 10 and f_{23} = 1

 - Renormalize and increment exponent (because of carry)
- ❖ Round towards +∞: Truncate result since negative
 - ♦ Truncated Result: -1.111111111111111111111 × 2-7
- ❖ Round towards $-\infty$: Increment since negative and $\mathbf{R} = \mathbf{1}$
- Round towards 0: Truncate always

Accuracy can be a Big Problem

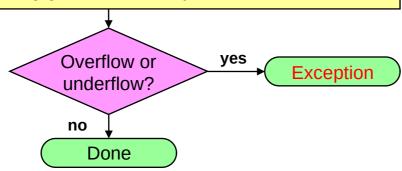
Value1	Value2	Value3	Value4	Sum
1.0E+30	-1.0E+30	9.5	-2.3	7.2
1.0E+30	9.5	-1.0E+30	-2.3	-2.3
1.0E+30	9.5	-2.3	-1.0E+30	0

- Adding double-precision floating-point numbers (Excel)
- Floating-Point addition is NOT associative
- Produces different sums for the same data values
- Rounding errors when the difference in exponent is large

Floating Point Addition / Subtraction

Start

- 1. Compare the exponents of the two numbers. Shift the smaller number to the right until its exponent would match the larger exponent.
- 2. Add / Subtract the significands according to the sign bits.
- 3. Normalize the sum, either shifting right and incrementing the exponent or shifting left and decrementing the exponent.
- 4. Round the significand to the appropriate number of bits, and renormalize if rounding generates a carry.



Shift significand right by $d = |E_x - E_y|$

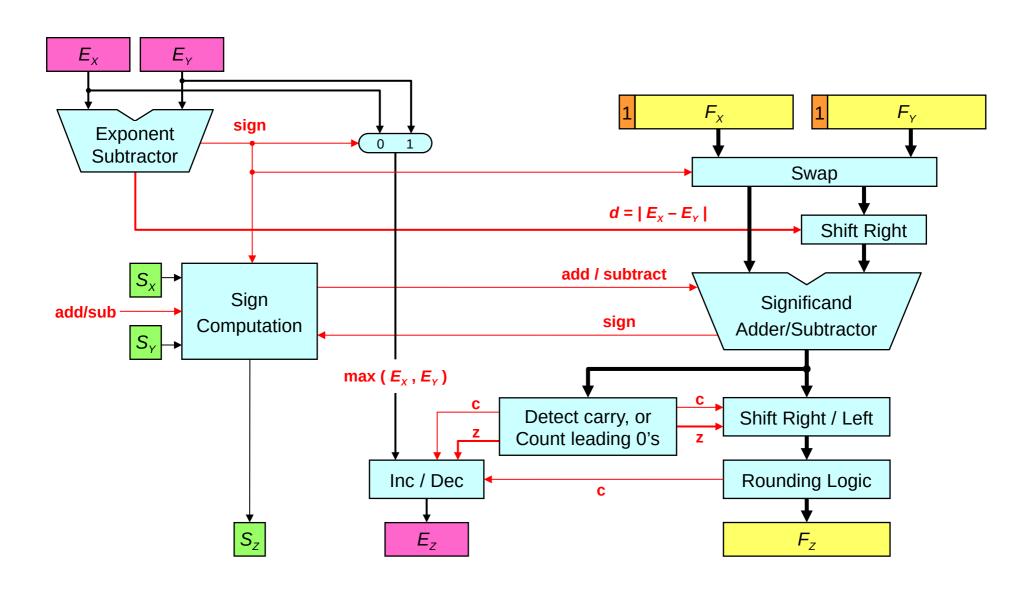
Add significands when signs of *X* and *Y* are identical, Subtract when different.

Convert negative result from 2's complement to sign-magnitude.

Normalization shifts right by 1 if there is a carry, or shifts left by the number of leading zeros in the case of subtraction.

Rounding either truncates fraction, or adds a 1 to least significant fraction bit.

Floating Point Adder Block Diagram





- Floating-Point Numbers
- The IEEE 754 Floating-Point Standard

- Floating-Point Comparison, Addition and Subtraction
- Floating-Point Multiplication
- MIPS Floating-Point Instructions and Examples

Floating Point Multiplication Example

Consider multiplying:

$$-1.110\ 1000\ 0100\ 0000\ 1010\ 0001_2\ \times\ 2^{-4}$$

- \times 1.100 0000 0001 0000 0000 0000₂ \times 2⁻²
- Unlike addition, we add the exponents of the operands
 - \diamond Result exponent value = (-4) + (-2) = -6
- Using the biased representation: $E_Z = E_X + E_Y Bias$

$$\Rightarrow$$
 $E_x = (-4) + 127 = 123$ (*Bias* = 127 for single precision)

$$\Rightarrow$$
 $E_{Y} = (-2) + 127 = 125$

$$\Rightarrow$$
 $E_7 = 123 + 125 - 127 = 121$ (exponent value = -6)

- Sign bit of product can be computed independently
- Sign bit of product = Sign XOR Sign Sign = 1 (negative)

Floating-Point Multiplication, cont'd

Now multiply the significands:

```
(Multiplicand)
```

1.11010000100000010100001

(Multiplier)

1.100000000010000000000000

111010000100000010100001

111010000100000010100001

- 1.11010000100000010100001
- 24 bits × 24 bits 48 bits (double number of bits)
- Multiplicand × 1 = Multiplicand (shifted left)

Floating-Point Multiplication, cont'd

Normalize Product

- $= -1.0101110001111110111111001100... \times 2^{-5}$
- * Round to Nearest Even: (keep only 23 fraction bits)

$$-1.010111000111111011111100 \mid 1 100... \times 2^{-5}$$

Round bit = 1, Sticky bit = 1, so increment fraction

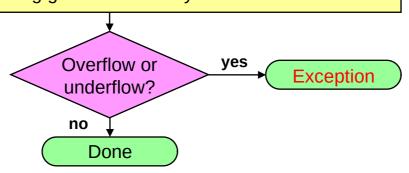
IEEE 754 Representation

10111101001011110001111110111111101

Floating Point Multiplication

Start

- 1. Add the biased exponents of the two numbers, subtracting the bias from the sum to get the new biased exponent
- 2. Multiply the significands. Set the result sign to positive if operands have same sign, and negative otherwise
- 3. Normalize the product if necessary, shifting its significand right and incrementing the exponent
- 4. Round the significand to the appropriate number of bits, and renormalize if rounding generates a carry



Biased Exponent Addition $E_z = E_x + E_y - Bias$

Result sign $S_Z = S_X \operatorname{\mathbf{xor}} S_Y \operatorname{\mathbf{can}}$ be computed independently

Since the operand significands $1.F_{\chi}$ and $1.F_{\gamma}$ are ≥ 1 and < 2, their product is ≥ 1 and < 4. To normalize product, we need to shift right at most by 1 bit and increment exponent

Rounding either truncates fraction, or adds a 1 to least significant fraction bit

Extra Bits to Maintain Precision

- Floating-point numbers are approximations for ...
 - ♦ Real numbers that they cannot represent
- Infinite real numbers exist between 1.0 and 2.0
 - ♦ However, exactly 2²³ fractions represented in Single Precision
 - ♦ Exactly 2⁵² fractions can be represented in Double Precision
- * Extra bits are generated in intermediate results when ...
 - ♦ Shifting and adding/subtracting a *p*-bit significand
 - ♦ Multiplying two *p*-bit significands (product is 2*p* bits)
- But when packing result fraction, extra bits are discarded
- Few extra bits are needed: guard, round, and sticky bits
- Minimize hardware but without compromising accuracy

Advantages of IEEE 754 Standard

- Used predominantly by the industry
- Encoding of exponent and fraction simplifies comparison
 - ♦ Integer comparator used to compare magnitude of FP numbers
- Includes special exceptional values: NaN and ±∞
 - ♦ Special rules are used such as:
 - 0/0 is NaN, sqrt(-1) is NaN, 1/0 is ∞ , and $1/\infty$ is 0
 - Computation may continue in the face of exceptional conditions
- Denormalized numbers to fill the gap
 - \diamond Between smallest normalized number $1.0 \times 2^{E_{min}}$ and zero
 - \diamond Denormalized numbers, values $0.F \times 2^{E_{min}}$, are closer to zero
 - Gradual underflow to zero

Floating Point Complexities

- Operations are somewhat more complicated
- In addition to overflow we can have underflow
- Accuracy can be a big problem
 - Extra bits to maintain precision: guard, round, and sticky
 - ♦ Four rounding modes
 - Division by zero yields Infinity
 - Zero divide by zero yields Not-a-Number
 - Other complexities
- Implementing the standard can be tricky
- Not using the standard can be even worse



Floating-Point Numbers

- The IEEE 754 Floating-Point Standard
- Floating-Point Comparison, Addition and Subtraction
- Floating-Point Multiplication
- MIPS Floating-Point Instructions and Examples

MIPS Floating Point Coprocessor

- Called Coprocessor 1 or the Floating Point Unit (FPU)
- 32 separate floating point registers: \$f0, \$f1, ..., \$f31
- FP registers are 32 bits for single precision numbers
- Even-odd register pair form a double precision register
- Use the even number for double precision registers
 \$f0, \$f2, \$f4, ..., \$f30 are used for double precision
- Separate FP instructions for single/double precision
 - ♦ Single precision: add.s, sub.s, mul.s, div.s (.s extension)
 - ♦ Double precision: add.d, sub.d, mul.d, div.d (.d extension)
- FP instructions are more complex than the integer ones
 - ♦ Take more cycles to execute

Floating-Point Arithmetic Instructions

Instruction		Meaning	Op ⁶	fmt ⁵	ft ⁵	fs ⁵	fd⁵	func ⁶
add.s	\$f5,\$f3,\$f4	\$f5 = \$f3 + \$f4	0x11	0x10	\$f4	\$f3	\$f5	0
sub.s	\$f5,\$f3,\$f4	\$f5 = \$f3 - \$f4	0x11	0x10	\$f4	\$f3	\$f5	1
mul.s	\$f5,\$f3,\$f4	\$f5 = \$f3 × \$f4	0x11	0x10	\$f4	\$f3	\$f5	2
div.s	\$f5,\$f3,\$f4	\$f5 = \$f3 / \$f4	0x11	0x10	\$f4	\$f3	\$f5	3
sqrt.s	\$f5,\$f3	\$f5 = sqrt(\$f3)	0x11	0x10	0	\$f3	\$f5	4
abs.s	\$f5,\$f3	\$f5 = abs(\$f3)	0x11	0x10	0	\$f3	\$f5	5
neg.s	\$f5,\$f3	f5 = -(f3)	0x11	0x10	0	\$f3	\$f5	7
add.d	\$f6,\$f2,\$f4	\$f6,7 = \$f2,3 + \$f4,5	0x11	0x11	\$f4	\$f2	\$f6	0
sub.d	\$f6,\$f2,\$f4	\$f6,7 = \$f2,3 - \$f4,5	0x11	0x11	\$f4	\$f2	\$f6	1
mul.d	\$f6,\$f2,\$f4	$$f6,7 = $f2,3 \times $f4,5$	0x11	0x11	\$f4	\$f2	\$f6	2
div.d	\$f6,\$f2,\$f4	\$f6,7 = \$f2,3 / \$f4,5	0x11	0x11	\$f4	\$f2	\$f6	3
sqrt.d	\$f6,\$f2	\$f6,7 = sqrt(\$f2,3)	0x11	0x11	0	\$f2	\$f6	4
abs.d	\$f6,\$f2	\$f6,7 = abs(\$f2,3)	0x11	0x11	0	\$f2	\$f6	5
neg.d	\$f6,\$f2	\$f6,7 = -(\$f2,3)	0x11	0x11	0	\$f2	\$f6	7

Floating-Point Load and Store

Separate floating-point load and store instructions

lwc1: load word coprocessor 1

1dc1: load double coprocessor 1

swc1: store word coprocessor 1

sdc1: store double coprocessor 1

General purpose register is used as the address register

Instruction	Meaning	Op ⁶	rs ⁵	ft ⁵	Immediate ¹⁶
lwc1 \$f2, 8(\$t0)	\$f2 7 ₄ Mem[\$t0+8]	0x31	\$t0	\$f2	8
swc1 \$f2, 8(\$t0)	\$f2 ♣4 Mem[\$t0+8]	0x39	\$t0	\$f2	8
ldc1 \$f2, 8(\$t0)	\$f2,3 % Mem[\$t0+8]	0x35	\$t0	\$f2	8
sdc1 \$f2 oating Point COE 301 / ICS 233 – Computer C	rgasization 3 A Muhama Madamsir + 19448]	0x3d	\$±0	\$f2	8

Data Movement Instructions

Moving data between general purpose and FP registers

mfc1: move from coprocessor 1 (to a general purpose register)

mtc1: move to coprocessor 1 (from a general purpose register)

Moving data between FP registers

mov.s: move single precision float

mov.d: move double precision float = even/odd pair of registers

Instruction			Meaning	0p ⁶	fmt ⁵	rt ⁵	fs ⁵	fd⁵	func
mfc1	\$t0,	\$f2	\$t0 = \$f2	0×11	0	\$t0	\$f2	0	0
mtc1	\$t0,	\$f2	\$f2 = \$t0	0×11	4	\$t0	\$f2	0	0
mov.s	\$f4,	\$f2	\$f4 = \$f2	0x11	0x10	0	\$f2	\$f4	6
mov.d	\$f4,	\$f2	\$f4,5 = \$f2,3	0x11	0x11	0	\$f2	\$f4	6

Convert Instructions

- Convert instruction: cvt.x.y
 - ♦ Convert the **source** format **y** into **destination** format **x**
- Supported Formats:
 - ♦ Single-precision float = .s
 - Double-precision float = .d
 - \diamond Signed integer word = .w (in a floating-point register)

Instruction	Meaning	0p ⁶	fmt ⁵		fs⁵	fd⁵	func
cvt.s.w \$f2,\$f4	\$f2 = W2S(\$f4)	0x11	0x14	0	\$f4	\$f2	0x20
cvt.s.d \$f2,\$f4	\$f2 = D2P(\$f4,5)	0x11	0x11	0	\$f4	\$f2	0x20
cvt.d.w \$f2,\$f4	\$f2,3 = W2D(\$f4)	0x11	0x14	0	\$f4	\$f2	0x21
cvt.d.s \$f2,\$f4	\$f2,3 = S2D(\$f4)	0x11	0x10	0	\$f4	\$f2	0x21
cvt.w.s \$f2,\$f4	\$f2 = S2W(\$f4)	0x11	0x10	0	\$f4	\$f2	0x24
cvt.w.d \$f2,\$f4	\$f2 = D2W(\$f4,5)	0x11	0x11	0	\$f4	\$f2	0x24

Floating-Point Compare and Branch

- Floating-Point unit has eight condition code cc flags
 - ♦ Set to 0 (false) or 1 (true) by any comparison instruction
- Three comparisons: eq (equal), It (less than), Ie (less or equal)
- Two branch instructions based on the condition flag

Instruction	Meaning	0p ⁶	fmt ⁵	ft ⁵	fs ⁵		func
c.eq.s cc \$f2,\$f4	cc = (\$f2 == \$f4)	0x11	0x10	\$f4	\$f2	СС	0x32
c.eq.d cc \$f2,\$f4	cc = (\$f2,3 == \$f4,5)	0x11	0x11	\$f4	\$f2	CC	0x32
c.lt.s cc \$f2,\$f4	cc = (\$f2 < \$f4)	0x11	0x10	\$f4	\$f2	CC	0x3c
c.lt.d cc \$f2,\$f4	cc = (\$f2,3 < \$f4,5)	0x11	0x11	\$f4	\$f2	СС	0x3c
c.le.s cc \$f2,\$f4	cc = (\$f2 <= \$f4)	0x11	0×10	\$f4	\$f2	СС	0x3e
c.le.d cc \$f2,\$f4	cc = (\$f2,3 <= \$f4,5)	0x11	0x11	\$f4	\$f2	CC	0x3e
bc1f cc Label	branch if (cc == 0)	0x11	8	cc,0	16-b	it 0	ffset
bc1t cc Label	branch if (cc == 1)	0x11	8	cc,1	16-b	it 0	ffset

Example 1: Area of a Circle

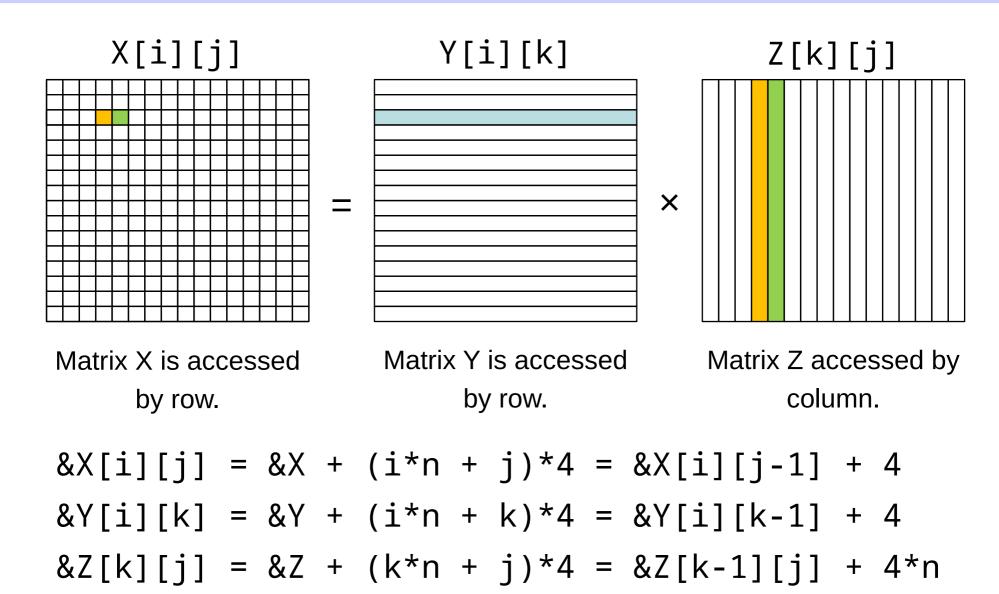
```
.data
 pi:
    .double 3.1415926535897924
 msg: .asciiz "Circle Area = "
.text
main:
 ldc1 $f2, pi # $f2,3 = pi
 li $v0, 7 # read double (radius)
         # $f0,1 = radius
 syscall
 mul.d$f12, $f0, $f0 # $f12,13 = radius*radius
 mul.d $f12, $f2, $f12 # $f12,13 = area
 la $a0, msg
 li $v0, 4 # print string (msg)
 syscall
 li $v0, 3
                # print double (area)
                 # print $f12,13
  syscall
```

Example 2: Matrix Multiplication

```
void mm (int n, float X[n][n], Y[n][n], Z[n][n]) {
  for (int i=0; i!=n; i=i+1) {
    for (int j=0; j!=n; j=j+1) {
       float sum = 0.0;
      for (int k=0; k!=n; k=k+1) {
        sum = sum + Y[i][k] * Z[k][j];
       X[i][j] = sum;
```

- Matrix size is passed in \$a0 = n
- ★ Matrix addresses in \$a1 = &X, \$a2 = &Y, and \$a3 = &Z
- What is the MIPS assembly code for the procedure?

Access Pattern for Matrix Multiply



Matrix Multiplication Procedure (1 of 3)

```
# arguments $a0=n, $a1=&X, $a2=&Y, $a3=&Z
mm: s11 $t0, $a0, 2 # $t0 = n*4 (row size)
      li $t1, 0 # $t1 = i = 0
# Outer for (i = . . .) loop starts here
L1: li $t2, 0 # $t2 = j = 0
# Middle for (j = . . . ) loop starts here
L2: 1i $t3, 0 # $t3 = k = 0
      move $t4, $a2 # $t4 = &Y[i][0]
      s11 + 5, +2, 2 + 5 = j*4
      addu $t5, $a3, $t5 # $t5 = &Z[0][j]
      mtc1 \$zero, \$f0  # \$f0 = sum = 0.0
```

Matrix Multiplication Procedure (2 of 3)

```
# Inner for (k = . . .) loop starts here
# $t3 = k, $t4 = &Y[i][k], $t5 = &Z[k][j]
L3: lwc1 $f1, 0($t4) # load $f1 = Y[i][k]
      lwc1 $f2, 0($t5) # load $f2 = Z[k][j]
      mul.s $f3, $f1, $f2 # $f3 = Y[i]
  [k]*Z[k][j]
      add.s
               $f0, $f0, $f3 # sum = sum + $f3
               $t3, $t3, 1  # k = k + 1
      addiu
      addiu $t4, $t4, 4 # $t4 = &Y[i][k]
      addu $t5, $t5, $t0 # $t5 = &Z[k][j]
               $t3, $a0, L3 # loop back if (k
      bne
```

Matrix Multiplication Procedure (3 of 3)

```
swc1 $f0, 0($a1) # store X[i][j] = sum
      addiu a1, a1, 4 # a1 = aX[i][j]
      addiu $t2, $t2, 1 # j = j + 1
      bne $t2, $a0, L2 # loop L2 if (j != n)
# End of middle for loop
      addu a2, a2, a2, a2, a2 = a2 = a2 = a2
      addiu $t1, $t1, 1 # i = i + 1
      bne $t1, $a0, L1 # loop L1 if (i != n)
# End of outer for loop
      jr $ra # return to caller
```