

Memory Hierarchy and Caches

COE 301 / ICS 233

Computer Organization

Dr. Muhamed Mudawar

College of Computer Sciences and Engineering
King Fahd University of Petroleum and Minerals

Presentation Outline

- ❖ **Random Access Memory and its Structure**
- ❖ Memory Hierarchy and the need for Cache Memory
- ❖ The Basics of Caches
- ❖ Cache Performance and Memory Stall Cycles
- ❖ Improving Cache Performance
- ❖ Multilevel Caches

Memory Technology

❖ Static RAM (SRAM)

- ◇ Used typically to implement Cache memory
- ◇ Requires 6 transistors per bit
- ◇ Low power to retain bit

❖ Dynamic RAM (DRAM)

- ◇ Used typically to implement Main Memory
- ◇ One transistor + capacitor per bit
- ◇ Must be re-written after being read
- ◇ Must be refreshed periodically
 - By reading and rewriting all the rows in the DRAM

Static RAM Storage Cell

- ❖ Static RAM (SRAM) Memory

- ❖ Typically used for **caches**

- ❖ Provides **fast access time**

- ❖ Cell Implementation:


 - ◇ 6-Transistor cell

 - ◇ Cross-coupled inverters store bit


 - ◇ Two pass transistors

 - ◇ Row decoder selects the word line

 - ◇ Pass transistors enable the cell to be read and written

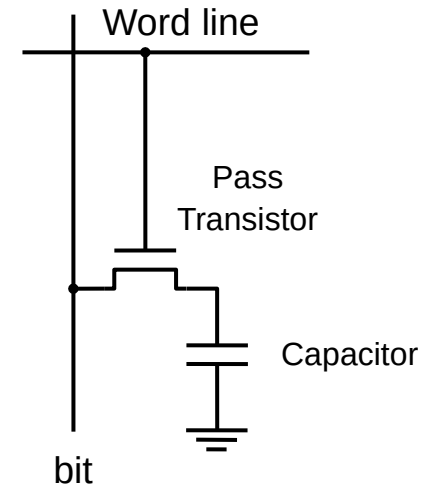


Typical SRAM cell



Dynamic RAM Storage Cell

- ❖ Dynamic RAM (DRAM): cheap, dense, but slower than SRAM
- ❖ Typical choice for **main memory**
- ❖ Cell Implementation:
 - ◇ 1-Transistor cell (pass transistor)
 - ◇ Trench capacitor (stores bit)
- ❖ Bit is stored as a **charge** on capacitor
- ❖ Must be **refreshed periodically**
 - ◇ Because of leakage of charge from tiny capacitor
- ❖ Refreshing for all memory rows
 - ◇ Reading each row and writing it back to restore the charge



Typical DRAM cell

Example of a Memory Chip

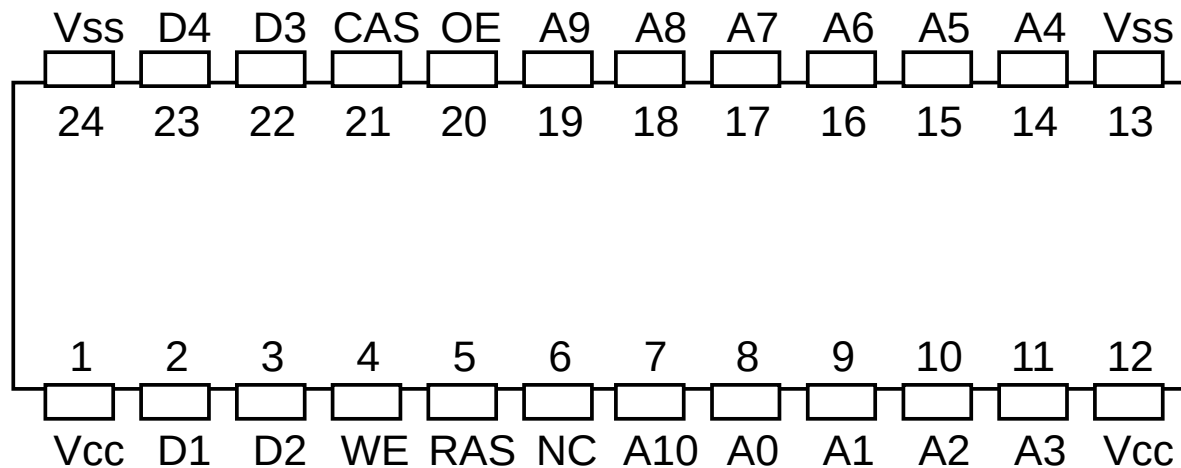
❖ 24-pin dual in-line package: $2^{22} \times 4\text{-bit} = 16\text{Mibit}$ memory

❖ 22-bit address is divided into

- ◇ 11-bit row address
- ◇ 11-bit column address
- ◇ Interleaved on same address lines

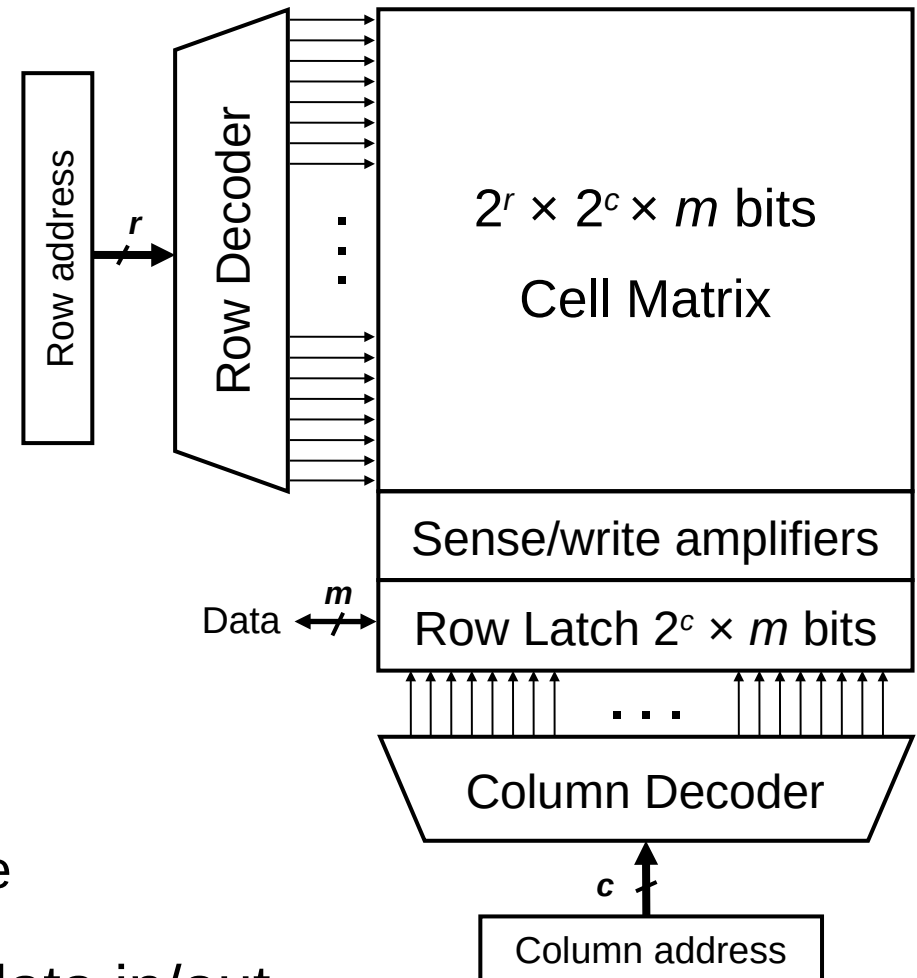
Legend

A_i	Address bit i
CAS	Column address strobe
D_j	Data bit j
NC	No connection
OE	Output enable
RAS	Row address strobe
WE	Write enable



Typical Memory Structure

- ❖ Row decoder
 - ◇ Select row to read/write
- ❖ Column decoder
 - ◇ Select column to read/write
- ❖ Cell Matrix
 - ◇ 2D array of tiny memory cells
- ❖ Sense/Write amplifiers
 - ◇ Sense & amplify data on read
 - ◇ Drive bit line with data in on write
- ❖ Same data lines are used for data in/out



DRAM Operation

❖ Row Access (RAS)

- ◇ Latch and decode row address to enable addressed row
- ◇ Small change in voltage detected by sense amplifiers
- ◇ Latch whole row of bits
- ◇ Sense amplifiers drive bit lines to recharge storage cells

❖ Column Access (CAS) read and write operation

- ◇ Latch and decode column address to select m bits
- ◇ $m = 4, 8, 16, \text{ or } 32$ bits depending on the DRAM package
- ◇ On read, send latched bits out to chip pins
- ◇ On write, charge storage cells to required value
- ◇ Can perform multiple column accesses to same row (burst mode)

Burst Mode Operation

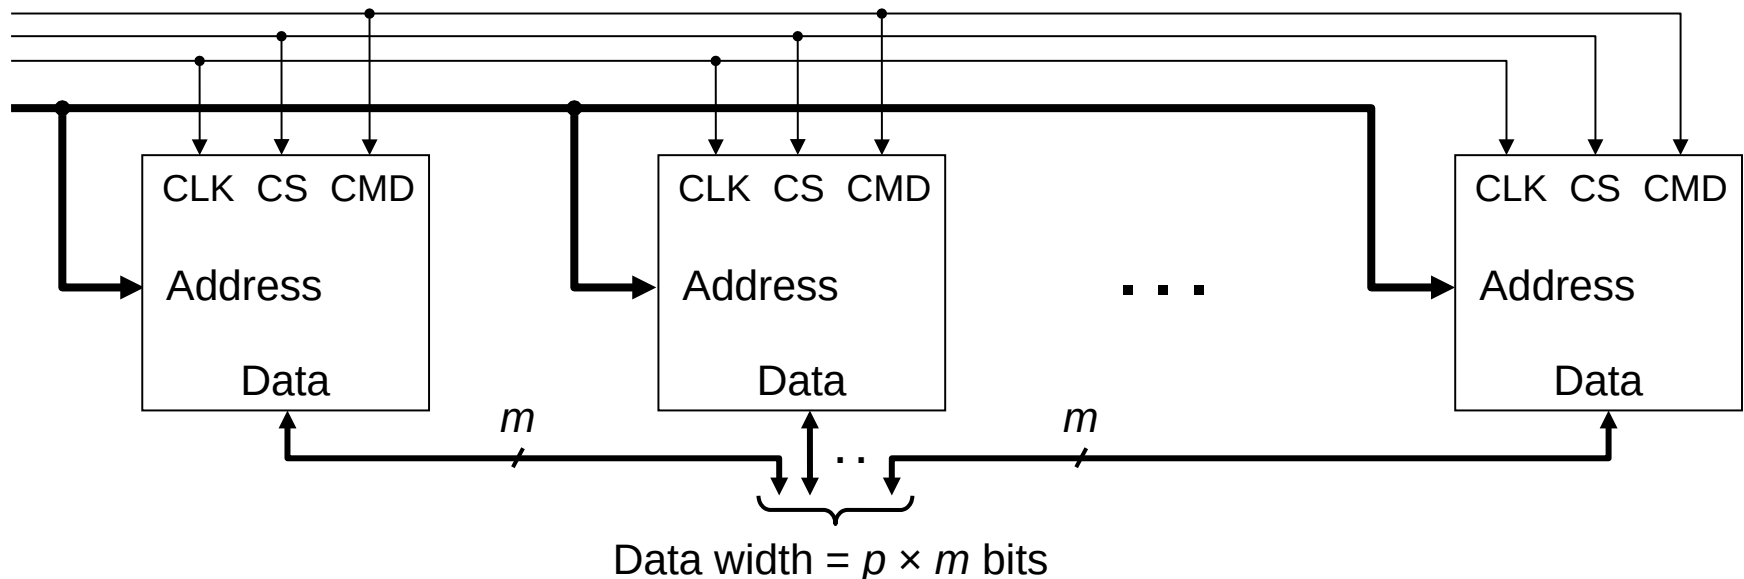
- ❖ Used for Block Transfer
 - ◇ Row address is latched and decoded
 - ◇ A read operation causes ALL cells in a selected row to be read
 - ◇ Selected row is latched internally inside the DRAM chip
 - ◇ Column address is latched and decoded
 - ◇ Selected column data is placed in the data output register
 - ◇ Column address is incremented automatically
 - ◇ Multiple data items are read depending on the block length
- ❖ Fast transfer of blocks between main memory and cache
- ❖ Fast transfer of pages between main memory and disk

SDRAM and DDR SDRAM

- ❖ SDRAM is **Synchronous** Dynamic RAM
 - ◇ Added clock to DRAM interface
- ❖ SDRAM is synchronous with the system clock
 - ◇ Older DRAM technologies were asynchronous
 - ◇ As system bus clock improved, SDRAM delivered higher performance than asynchronous DRAM
- ❖ DDR is **Double Data Rate** SDRAM
 - ◇ Like SDRAM, DDR is synchronous with the system clock, but the difference is that DDR reads data on both the rising and falling edges of the clock signal


Memory Modules

- ❖ Memory Rank: Set of DRAM chips accessed in parallel
 - ◇ Same Chip Select (CS) and Command (CMD)
 - ◇ Same address, but different data lines
 - ◇ Increases memory capacity and bandwidth
 - ◇ Example: 64-bit data bus using 4 × 16-bit DRAM chips



Trends in DRAM

Year	Memory Standard	Chip Capacity (Mibit)	Bus Clock (MHz)	Data Rate (MT/s)	Peak Bandwidth (MB/s)	Total latency to a new row / column
1996	SDRAM	64-128	100-166	100-166	800-1333	60 ns
2000	DDR	256-512	100-200	200-400	1600-3200	55 ns
2004	DDR2	512-2048	200-400	400-800	3200-6400	50 ns
2010	DDR3	2048-8192	400-800	800-1600	6400-12800	40 ns
2014	DDR4	8192-32768	800-1600	1600-3200	12800-25600	35 ns

- ❖ Memory chip capacity: 1 Mibit = 2^{20} bits, 1 Gibit = 2^{30} bits
- ❖ Data Rate = Millions of Transfers per second (MT/s)
- ❖ Data Rate = $2 \times$ Bus Clock for DDR, DDR2, DDR3, DDR4
- ❖ 1 Transfer = 8 bytes of data  Bandwidth = MT/s \times 8 bytes

Memory Latency versus Bandwidth

❖ Memory Latency


- ◇ Elapsed time between sending address and receiving data
- ◇ Measured in nanoseconds
- ◇ The total latency to a new row/column is the time between opening a new row of memory and accessing a column within it.
- ◇ Reduced from 60 ns to 35 ns (between 1996 and 2016)
- ◇ Improvement in memory latency is less than 2X (1996 to 2016)

❖ Memory Bandwidth

- ◇ Rate at which data is transferred between memory and CPU
- ◇ Bandwidth is measured as millions of Bytes per second
- ◇ Increased from 800 to 25600 MBytes/sec (between 1996 and 2016)
- ◇ Improvement in memory bandwidth is 32X (1996 to 2016)

DRAM Refresh Cycles


- ❖ Refresh cycle is about tens of milliseconds
- ❖ Refreshing is done for the entire memory
- ❖ Each row is read and written back to restore the charge
- ❖ Some of the memory bandwidth is lost to refresh cycles



Next . . .

- ❖ Random Access Memory and its Structure
- ❖ **Memory Hierarchy and the need for Cache Memory**
- ❖ The Basics of Caches
- ❖ Cache Performance and Memory Stall Cycles
- ❖ Improving Cache Performance
- ❖ Multilevel Caches

Processor-Memory Performance Gap



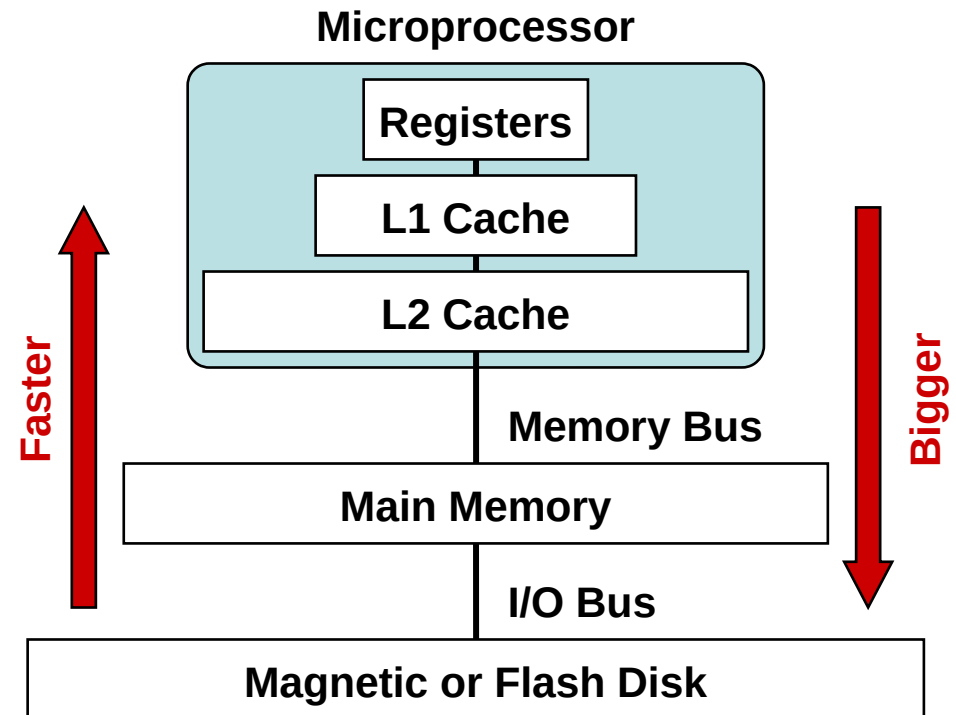
- ❖ 1980 – No cache in microprocessor
- ❖ 1995 – Two-level cache on microprocessor

The Need for Cache Memory

- ❖ Widening speed gap between CPU and main memory
 - ◇ Processor operation takes less than 1 ns
 - ◇ Main memory requires more than 50 ns to access
- ❖ Each instruction involves at least one memory access
 - ◇ One memory access to fetch the instruction
 - ◇ A second memory access for load and store instructions
- ❖ Memory bandwidth limits the instruction execution rate
- ❖ Cache memory can help bridge the CPU-memory gap
- ❖ Cache memory is small in size but fast

Typical Memory Hierarchy

- ❖ Registers are at the top of the hierarchy
 - ◇ Typical size < 1 KB
 - ◇ Access time < 0.5 ns
- ❖ Level 1 Cache (8 – 64 KiB)
 - ◇ Access time: 1 ns
- ❖ L2 Cache (1 MiB – 8 MiB)
 - ◇ Access time: 3 – 10 ns
- ❖ Main Memory (8 – 32 GiB)
 - ◇ Access time: 40 – 50 ns
- ❖ Disk Storage (> 200 GB)
 - ◇ Access time: 5 – 10 ms




Principle of Locality of Reference

- ❖ Programs access small portion of their address space
 - ◇ At any time, only a small set of instructions & data is needed
- ❖ **Temporal Locality** (in time)
 - ◇ If an item is accessed, probably it will be accessed again soon
 - ◇ Same loop instructions are fetched each iteration
 - ◇ Same procedure may be called and executed many times
- ❖ **Spatial Locality** (in space)
 - ◇ Tendency to access contiguous instructions/data in memory
 - ◇ Sequential execution of Instructions
 - ◇ Traversing arrays element by element

What is a Cache Memory ?

- ❖ Small and fast (SRAM) memory technology
 - ◇ Stores the subset of instructions & data currently being accessed
- ❖ Used to reduce average access time to memory
- ❖ Caches exploit **temporal locality** by ...
 - ◇ Keeping recently accessed data closer to the processor
- ❖ Caches exploit **spatial locality** by ...
 - ◇ Moving blocks consisting of multiple contiguous words
- ❖ Goal is to achieve
 - ◇ **Fast speed** of cache memory access
 - ◇ Balance the **cost** of the memory system

Cache Memories in the Datapath



Almost Everything is a Cache !

- ❖ In computer architecture, almost everything is a cache!
- ❖ Registers: a **cache on variables** – software managed
- ❖ First-level cache: a **cache on second-level cache**
- ❖ Second-level cache: a **cache on memory** (or L3 cache)
- ❖ Memory: a **cache on hard disk**
 - ◇ Stores recent programs and their data
 - ◇ Hard disk can be viewed as an extension to main memory
- ❖ Branch target and prediction buffer
 - ◇ **Cache on branch target and prediction** information

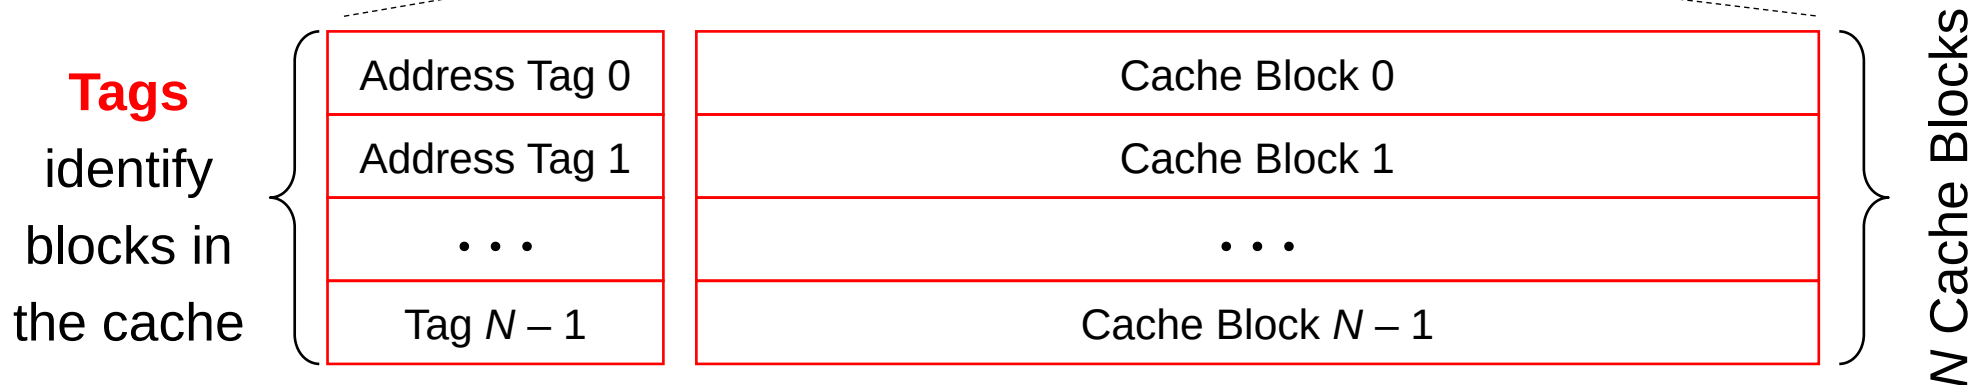

Next . . .

- ❖ Random Access Memory and its Structure
- ❖ Memory Hierarchy and the need for Cache Memory
- ❖ **The Basics of Caches**
- ❖ Cache Performance and Memory Stall Cycles
- ❖ Improving Cache Performance
- ❖ Multilevel Caches


Four Basic Questions on Caches

- ❖ Q1: Where can a block be placed in a cache?
 - ◇ **Block placement**
 - ◇ Direct Mapped, Set Associative, Fully Associative
- ❖ Q2: How is a block found in a cache?
 - ◇ **Block identification**
 - ◇ Block address, tag, index
- ❖ Q3: Which block should be replaced on a cache miss?
 - ◇ **Block replacement**
 - ◇ FIFO, Random, LRU
- ❖ Q4: What happens on a write?
 - ◇ **Write strategy**
 - ◇ Write Back or Write Through cache (with Write Buffer)

Inside a Cache Memory



❖ Cache Block (or Cache Line)


- ❖ Unit of data transfer between main memory and a cache
- ❖ Large block size  Less tag overhead + Burst transfer from DRAM
- ❖ Typically, cache block size = 64 bytes in recent caches

Block Placement: Direct Mapped

❖ **Block:** unit of data transfer between cache and memory

❖ **Direct Mapped Cache:**


◇ A block can be placed in exactly one location in the cache



Direct-Mapped Cache


- ❖ A memory address is divided into
 - ❖ **Block address**: identifies block in memory
 - ❖ **Block offset**: to access bytes within a block
- ❖ A block address is further divided into
 - ❖ **Index**: used for direct cache access
 - ❖ **Tag**: most-significant bits of block address

$$\text{Index} = \text{Block Address} \bmod \text{Cache Blocks}$$
- ❖ Tag must be stored also inside cache
 - ❖ For block identification
- ❖ A **valid bit** is also required to indicate
 - ❖ Whether a cache block is valid or not



Direct Mapped Cache – cont'd

- ❖ **Cache hit**: block is stored inside cache
 - ◇ Index is used to access cache block
 - ◇ Address tag is compared against stored tag
 - ◇ If equal and cache block is valid then **hit**
 - ◇ Otherwise: **cache miss**
- ❖ If number of cache blocks is 2^n
 - ◇ n bits are used for the cache index
- ❖ If number of bytes in a block is 2^b
 - ◇ b bits are used for the block offset
- ❖ If 32 bits are used for an address
 - ◇ $32 - n - b$ bits are used for the tag
- ❖ Cache data size = 2^{n+b} bytes




Mapping an Address to a Cache Block

❖ Example

- ❖ Consider a direct-mapped cache with 256 blocks
- ❖ Block size = 16 bytes
- ❖ Compute tag, index, and byte offset of address: 0x01FFF8AC

❖ Solution

- ❖ 32-bit address is divided into:



- 4-bit byte offset field, because block size = $2^4 = 16$ bytes
 - 8-bit cache index, because there are $2^8 = 256$ blocks in cache
 - 20-bit tag field
- ❖ Byte offset = 0xC = 12 (least significant 4 bits of address)
 - ❖ Cache index = 0x8A = 138 (next lower 8 bits of address)
 - ❖ Tag = 0x01FFF (upper 20 bits of address)

Example on Cache Placement & Misses

- ❖ Consider a small direct-mapped cache with 32 blocks
 - ◇ Cache is initially empty, Block size = 16 bytes
 - ◇ The following memory addresses (in decimal) are referenced:
1000, 1004, 1008, 2548, 2552, 2556.
 - ◇ Map addresses to cache blocks and indicate whether hit or miss

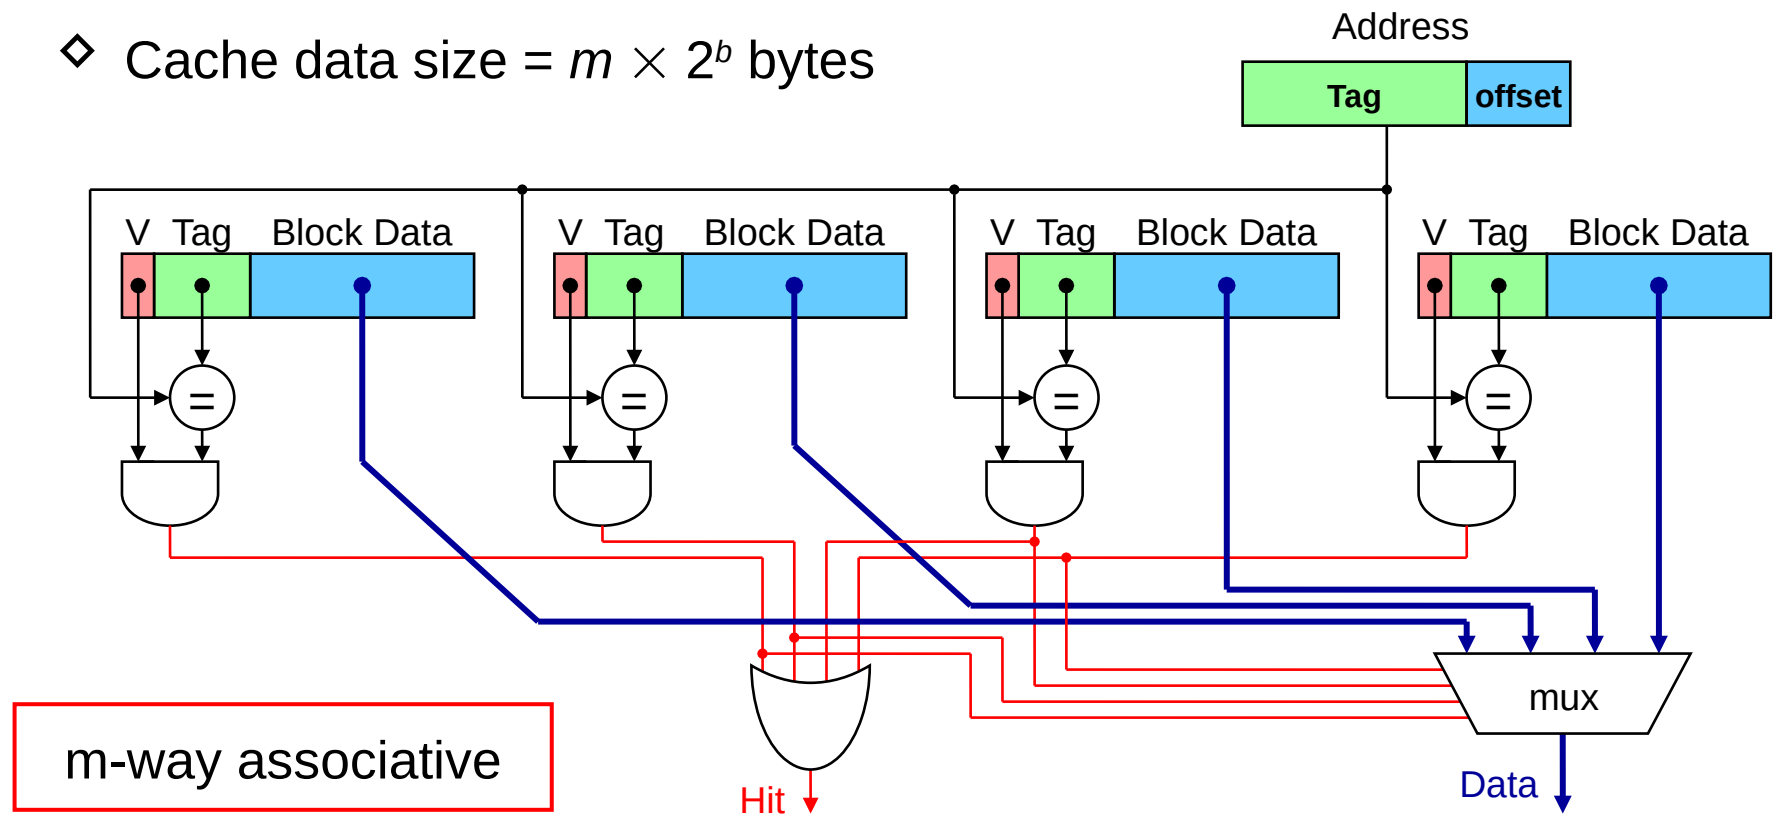
❖ Solution:



- ◇ 1000 = 0x3E8 cache index = 0x1E Miss (first access)
- ◇ 1004 = 0x3EC cache index = 0x1E Hit
- ◇ 1008 = 0x3F0 cache index = 0x1F Miss (first access)
- ◇ 2548 = 0x9F4 cache index = 0x1F Miss (different tag)
- ◇ 2552 = 0x9F8 cache index = 0x1F Hit
- ◇ 2556 = 0x9FC cache index = 0x1F Hit

Fully Associative Cache

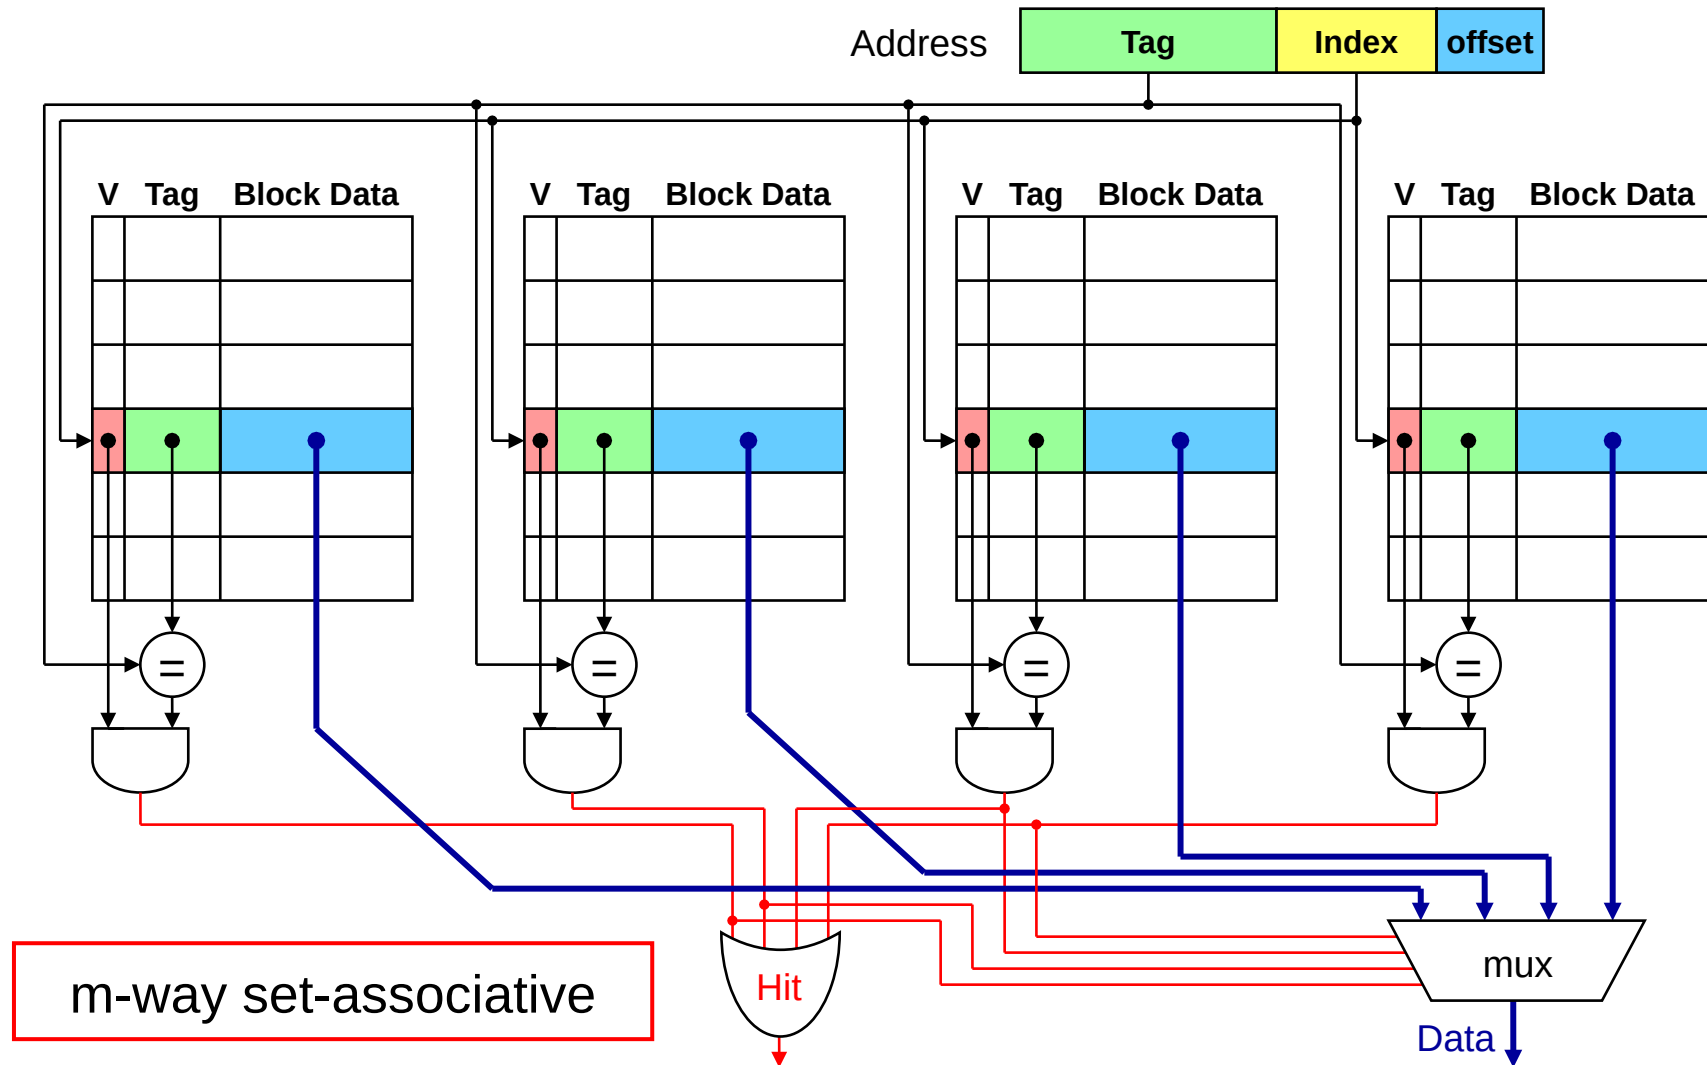
- ❖ A block can be placed **anywhere** in cache \Rightarrow **no indexing**
- ❖ If m blocks exist then
 - ◇ m comparators are needed to match *tag*
 - ◇ Cache data size = $m \times 2^b$ bytes



Set-Associative Cache

- ❖ A **set** is a group of blocks that can be indexed
- ❖ A block is first mapped onto a set
 - ◇ *Set index = Block address mod Number of sets in cache*
- ❖ If there are m blocks in a set (**m -way set associative**) then
 - ◇ m tags are checked in parallel using m comparators
- ❖ If 2^n sets exist then **set index** consists of n bits
- ❖ Cache data size = $m \times 2^{n+b}$ bytes (with 2^b bytes per block)
 - ◇ Without counting tags and valid bits
- ❖ A direct-mapped cache has one block per set ($m = 1$)
- ❖ A fully-associative cache has one set ($2^n = 1$ or $n = 0$)

Set-Associative Cache Diagram



Write Policy

❖ Write Through:

- ◇ Writes update cache and lower-level memory
- ◇ Cache control bit: only a **Valid** bit is needed
- ◇ Memory always has **latest data**, which simplifies data coherency
- ◇ Can always discard cached data when a block is replaced

❖ Write Back:

- ◇ Writes update cache only
- ◇ Cache control bits: **Valid** and **Modified** bits are required
- ◇ **Modified** cached data is **written back** to memory **when replaced**
- ◇ Multiple writes to a cache block require only one write to memory
- ◇ Uses **less memory bandwidth** than write-through and less power
- ◇ However, more complex to implement than write through

What Happens on a Cache Miss?

- ❖ Cache sends a **miss signal** to **stall** the processor
- ❖ Decide which cache block to allocate/replace
 - ◇ One choice only when the cache is directly mapped
 - ◇ Multiple choices for set-associative or fully-associative cache
- ❖ Transfer the block from lower level memory to this cache
 - ◇ Set the valid bit and the tag field from the upper address bits
- ❖ If block to be replaced is **modified** then write it back
 - ◇ Modified block is written back to memory
 - ◇ Otherwise, block to be replaced can be simply discarded
- ❖ Restart the instruction that caused the cache miss
- ❖ **Miss Penalty:** clock cycles to process a cache miss

Replacement Policy

- ❖ Which block to be replaced on a cache miss?
- ❖ No selection alternatives for direct-mapped caches
- ❖ m blocks per set to choose from for associative caches
- ❖ **Random replacement**
 - ◇ Candidate blocks are randomly selected
 - ◇ **One counter for all sets** (0 to $m - 1$): incremented on every cycle
 - ◇ On a cache miss replace block specified by counter
- ❖ **First In First Out (FIFO) replacement**
 - ◇ Replace oldest block in set
 - ◇ **One counter per set** (0 to $m - 1$): specifies **oldest block** to replace
 - ◇ Counter is incremented on a cache miss

Replacement Policy – cont'd

❖ Least Recently Used (LRU)

- ◇ Replace block that has been **unused for the longest time**
- ◇ Order blocks within a set from least to most recently used
- ◇ Update ordering of blocks on each cache hit
- ◇ With m blocks per set, there are **$m!$** possible permutations

❖ Pure LRU **is too costly** to implement when $m > 2$

- ◇ $m = 2$, there are 2 permutations only (a single bit is needed)
- ◇ $m = 4$, there are $4! = 24$ possible permutations
- ◇ LRU approximation is used in practice

❖ For large $m > 4$,

Random replacement can be as effective as LRU

Next . . .

- ❖ Random Access Memory and its Structure
- ❖ Memory Hierarchy and the need for Cache Memory
- ❖ The Basics of Caches
- ❖ **Cache Performance and Memory Stall Cycles**
- ❖ Improving Cache Performance
- ❖ Multilevel Caches

Hit Rate and Miss Rate

- ❖ Hit Rate = Hits / (Hits + Misses)
- ❖ Miss Rate = Misses / (Hits + Misses)
- ❖ I-Cache Miss Rate = Miss rate in the Instruction Cache
- ❖ D-Cache Miss Rate = Miss rate in the Data Cache
- ❖ Example:
 - ◇ Out of 1000 instructions fetched, 150 missed in the I-Cache
 - ◇ 25% are load-store instructions, 50 missed in the D-Cache
 - ◇ What are the I-cache and D-cache miss rates?
- ❖ I-Cache Miss Rate = $150 / 1000 = 15\%$
- ❖ D-Cache Miss Rate = $50 / (25\% \times 1000) = 50 / 250 = 20\%$

Memory Stall Cycles

- ❖ The processor **stalls** on a Cache miss
 - ◇ When fetching instructions from the **Instruction Cache (I-cache)**
 - ◇ When loading or storing data into the **Data Cache (D-cache)**

Memory stall cycles = Combined Misses \times Miss Penalty

- ❖ **Miss Penalty:** clock cycles to process a cache miss

Combined Misses = I-Cache Misses + D-Cache Misses

I-Cache Misses = I-Count \times I-Cache Miss Rate

D-Cache Misses = LS-Count \times D-Cache Miss Rate

LS-Count (Load & Store) = I-Count \times LS Frequency

- ❖ Cache misses are often reported per thousand instructions

Memory Stall Cycles Per Instruction

- ❖ **Memory Stall Cycles Per Instruction =**
Combined Misses Per Instruction \times Miss Penalty
- ❖ Miss Penalty is assumed equal for I-cache & D-cache
- ❖ Miss Penalty is assumed equal for Load and Store
- ❖ **Combined Misses Per Instruction =**
I-Cache Miss Rate + LS Frequency \times D-Cache Miss Rate
- ❖ **Therefore, Memory Stall Cycles Per Instruction =**
I-Cache Miss Rate \times Miss Penalty +
LS Frequency \times D-Cache Miss Rate \times Miss Penalty

Example on Memory Stall Cycles

- ❖ Consider a program with the given characteristics
 - ◇ Instruction count (**I-Count**) = 10^6 instructions
 - ◇ 30% of instructions are loads and stores
 - ◇ D-cache miss rate is 5% and I-cache miss rate is 1%
 - ◇ Miss penalty is 100 clock cycles for instruction and data caches
 - ◇ Compute combined misses per instruction and memory stall cycles
- ❖ **Combined misses per instruction in I-Cache and D-Cache**
 - ◇ $1\% + 30\% \times 5\% = 0.025$ combined misses per instruction
 - ◇ Equal to 25 misses per 1000 instructions
- ❖ **Memory stall cycles**
 - ◇ 0.025×100 (miss penalty) = 2.5 stall cycles per instruction
 - ◇ Total memory stall cycles = $10^6 \times 2.5 = 2,500,000$

CPU Time with Memory Stall Cycles

$$\text{CPU Time} = \text{I-Count} \times \text{CPI}_{\text{MemoryStalls}} \times \text{Clock Cycle}$$

$$\text{CPI}_{\text{MemoryStalls}} = \text{CPI}_{\text{PerfectCache}} + \text{Mem Stalls per Instruction}$$

- ❖ $\text{CPI}_{\text{PerfectCache}}$ = CPI for ideal cache (no cache misses)
- ❖ $\text{CPI}_{\text{MemoryStalls}}$ = CPI in the presence of memory stalls
- ❖ Memory stall cycles increase the CPI

Example on CPI with Memory Stalls

- ❖ A processor has CPI of 1.5 without any memory stalls
 - ◇ Cache miss rate is 2% for instruction and 5% for data
 - ◇ 20% of instructions are loads and stores
 - ◇ Cache miss penalty is 100 clock cycles for I-cache and D-cache
- ❖ What is the impact on the CPI?

❖ **Answer:**

$$\text{Mem Stalls per Instruction} = \underbrace{0.02 \times 100}_{\text{Instruction}} + \underbrace{0.2 \times 0.05 \times 100}_{\text{data}} = 3$$

$$\text{CPI}_{\text{MemoryStalls}} = 1.5 + 3 = 4.5 \text{ cycles per instruction}$$

$$\text{CPI}_{\text{MemoryStalls}} / \text{CPI}_{\text{PerfectCache}} = 4.5 / 1.5 = 3$$

Processor is **3 times slower** due to memory stall cycles

Average Memory Access Time

❖ Average Memory Access Time (AMAT)

$$\text{AMAT} = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$$

❖ Time to access a cache for both hits and misses

❖ Example: Find the AMAT for a cache with

- ❖ Cache access time (Hit time) of 1 cycle = 2 ns
- ❖ Miss penalty of 20 clock cycles
- ❖ Miss rate of 0.05 per access

❖ Solution:

$$\text{AMAT} = 1 + 0.05 \times 20 = 2 \text{ cycles} = 4 \text{ ns}$$

Without the cache, AMAT will be equal to Miss penalty = 20 cycles

Next . . .

- ❖ Random Access Memory and its Structure
- ❖ Memory Hierarchy and the need for Cache Memory
- ❖ The Basics of Caches
- ❖ Cache Performance and Memory Stall Cycles
- ❖ **Improving Cache Performance**
- ❖ **Multilevel Caches**

Improving Cache Performance

❖ **Average Memory Access Time (AMAT)**

$$\text{AMAT} = \text{Hit time} + \text{Miss rate} * \text{Miss penalty}$$

❖ Used as a framework for optimizations

❖ **Reduce the Hit time**

- ◇ Small and simple caches

❖ **Reduce the Miss Rate**

- ◇ Larger cache size, higher associativity, and larger block size

❖ **Reduce the Miss Penalty**

- ◇ Multilevel caches

Small and Simple Caches

- ❖ **Hit time is critical:** affects the processor clock cycle
 - ◇ Fast clock rate demands small and simple L1 cache designs
- ❖ Small cache reduces the indexing time and hit time
 - ◇ Indexing a cache represents a time consuming portion
 - ◇ Tag comparison also adds to this hit time
- ❖ Direct-mapped overlaps tag check with data transfer
 - ◇ Associative cache uses additional mux and increases hit time
- ❖ Size of L1 caches has not increased much
 - ◇ L1 caches are the same size on Alpha 21264 and 21364
 - ◇ Same also on UltraSparc II and III, AMD K6 and Athlon
 - ◇ Reduced from 16 KB in Pentium III to 8 KB in Pentium 4!

Classifying Misses – Three Cs


- ❖ Conditions under which misses occur
 - ❖ **Compulsory**: program starts with no block in cache
 - ◇ Also called **cold start misses**
 - ◇ Misses that would occur even if a cache has infinite size
 - ❖ **Capacity**: misses happen because cache size is finite
 - ◇ Blocks are replaced and then later retrieved
 - ◇ Misses that would occur in a fully associative cache of a finite size
 - ❖ **Conflict**: misses happen because of limited associativity
 - ◇ Limited number of blocks per set
 - ◇ Non-optimal replacement algorithm

Classifying Misses – cont'd

Compulsory misses are independent of cache size

Very small for long-running programs

Miss Rate




Larger Size and Higher Associativity

- ❖ Increasing cache size reduces capacity misses
- ❖ It also reduces conflict misses
 - ◇ Larger cache size spreads out references to more blocks
- ❖ Drawbacks: longer hit time and higher cost
- ❖ Larger caches are especially popular as 2nd level caches
- ❖ Higher associativity also improves miss rates
 - ◇ Eight-way set associative is as effective as a fully associative

Larger Block Size

- ❖ Simplest way to reduce miss rate is to increase block size
- ❖ However, it increases conflict misses if cache is small



Multilevel Caches

❖ Top level cache should be kept small to

◇ Keep pace with processor speed

❖ Adding another cache level

◇ Can reduce the memory gap

◇ Can reduce memory bus loading

❖ Local miss rate


◇ Number of misses in a cache / Memory accesses to this cache

◇ Miss Rate_{L_1} for L1 cache, and Miss Rate_{L_2} for L2 cache

❖ Global miss rate

Number of misses in a cache / Memory accesses generated by CPU

Miss Rate_{L_1} for L1 cache, and $\text{Miss Rate}_{L_1} \times \text{Miss Rate}_{L_2}$ for L2 cache

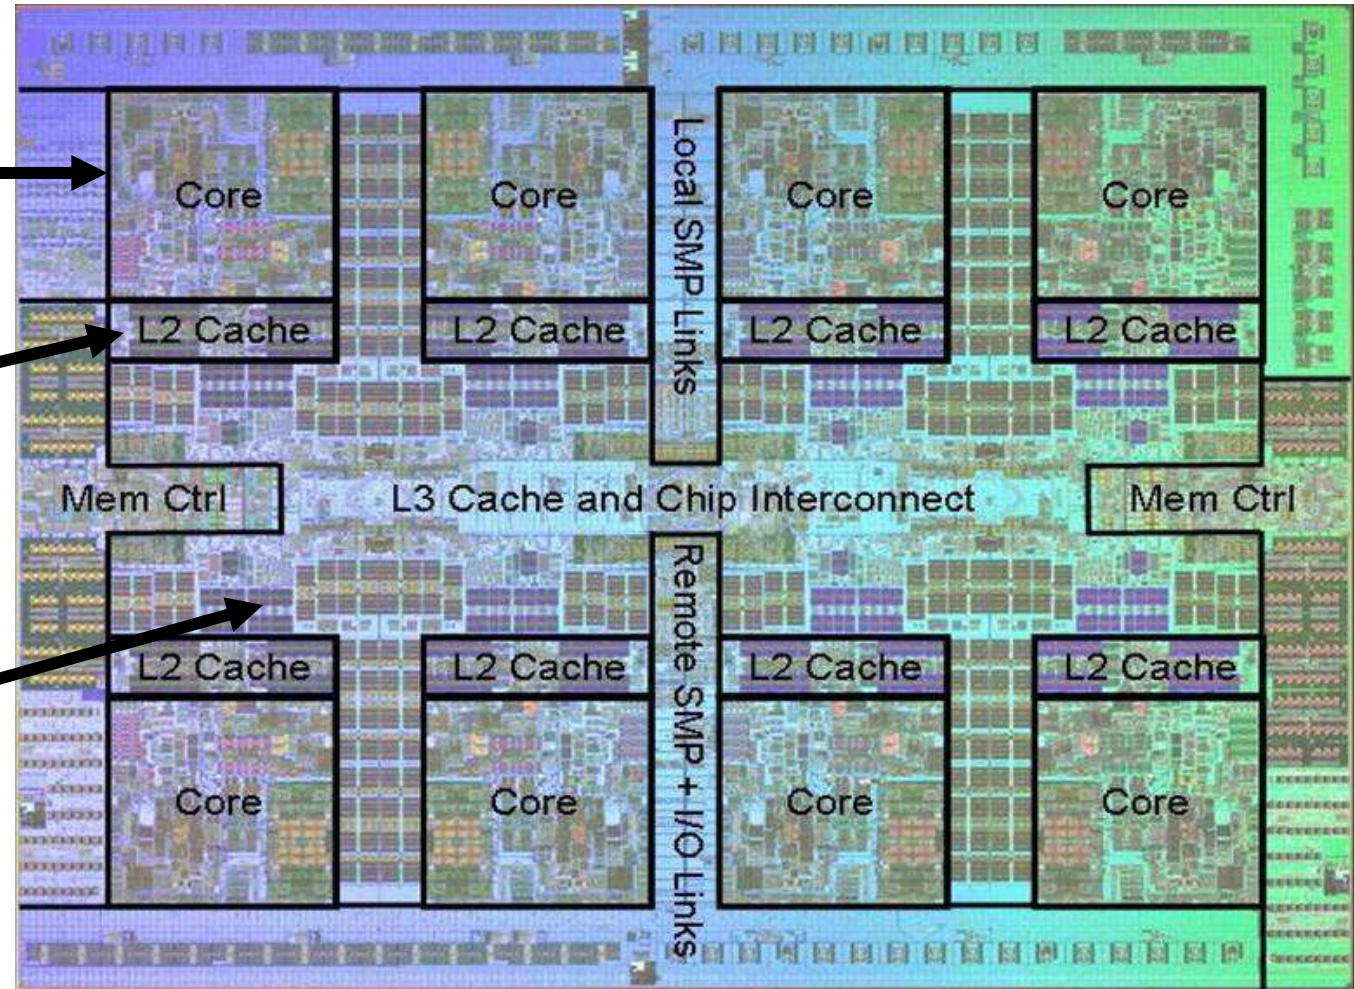


Power 7 On-Chip Caches [IBM 2010]

32KB I-Cache/core
32KB D-Cache/core
3-cycle latency

256KB Unified
L2 Cache/core
8-cycle latency

32MB Unified
Shared L3 Cache
Embedded DRAM
25-cycle latency
to local slice



Multilevel Cache Policies

❖ **Multilevel Inclusion**

- ◇ L1 cache data is always present in L2 cache
- ◇ A miss in L1, but a hit in L2 copies block from L2 to L1
- ◇ A miss in L1 and L2 brings a block into L1 and L2
- ◇ A write in L1 causes data to be written in L1 and L2
- ◇ Typically, write-through policy is used from L1 to L2
- ◇ Typically, write-back policy is used from L2 to main memory
 - To reduce traffic on the memory bus
- ◇ A replacement or invalidation in L2 must be propagated to L1

Multilevel Cache Policies – cont'd

❖ **Multilevel exclusion**

- ◇ L1 data is never found in L2 cache – Prevents wasting space
- ◇ Cache miss in L1, but a hit in L2 results in a swap of blocks
- ◇ Cache miss in both L1 and L2 brings the block into L1 only
- ◇ Block replaced in L1 is moved into L2
- ◇ Example: AMD Athlon

❖ **Same or different block size in L1 and L2 caches**

- ◇ Choosing a larger block size in L2 can improve performance
- ◇ However different block sizes complicates implementation
- ◇ Pentium 4 has 64-byte blocks in L1 and 128-byte blocks in L2